

Name: Hanish Sai Rohit Ph no: 8332082623 Email id : hanishsidhu@gmail.com Email id : hanishrohit@gmail.com

```
In [296]: from sklearn.datasets import load_boston
boston = load_boston()
```

```
In [297]: import pandas as pd
bos = pd.DataFrame(boston.data)
print(bos.head())
```

	0	1	2	3	4	5	6	7	8	9	10	\
0	0.00632	18.0	2.31	0.0	0.538	6.575	65.2	4.0900	1.0	296.0	15.3	
1	0.02731	0.0	7.07	0.0	0.469	6.421	78.9	4.9671	2.0	242.0	17.8	
2	0.02729	0.0	7.07	0.0	0.469	7.185	61.1	4.9671	2.0	242.0	17.8	
3	0.03237	0.0	2.18	0.0	0.458	6.998	45.8	6.0622	3.0	222.0	18.7	
4	0.06905	0.0	2.18	0.0	0.458	7.147	54.2	6.0622	3.0	222.0	18.7	

	11	12
0	396.90	4.98
1	396.90	9.14
2	392.83	4.03
3	394.63	2.94
4	396.90	5.33

```
In [298]: bos['PRICE'] = boston.target
```

```
X = bos.drop('PRICE', axis = 1)
Y = bos['PRICE']
```

```
In [299]: import sklearn.cross_validation
X_train, X_test, Y_train, Y_test = sklearn.cross_validation.train_test_split(X, Y, test_size = 0.33, random_state = 5)
print(X_train.shape)
print(X_test.shape)
print(Y_train.shape)
print(Y_test.shape)
```

```
(339, 13)
(167, 13)
(339,)
(167,)
```

```
In [300]: from sklearn.preprocessing import StandardScaler
scaler = StandardScaler()
scaler.fit(X_train)
X_train = scaler.transform(X_train)
X_test = scaler.transform(X_test)
```

```
In [301]: X_train = pd.DataFrame(X_train)
X_test = pd.DataFrame(X_test)
```

Scikit-learn Implementation

```
In [302]: # code source: https://medium.com/@haydar\_ai/learning-data-science-day-9-linear-regression-on-boston-housing-dataset-cd62a80775ef
from sklearn.linear_model import LinearRegression
import matplotlib.pyplot as plt

lm = LinearRegression()
lm.fit(X_train, Y_train)

pred_sklearn = lm.predict(X_test)
```

Stochastic Gradient Descent Implementation

```

In [354]: import numpy as np

class Stochastic_Gradient_Descent:

    # a function to calculate partial derivative of weights
    def partial_derivative_of_weights(self, weight, intercept, train_X, train_Y):
        xw = np.matmul(train_X, weight)
        xwi = xw + intercept
        result = np.subtract(xwi, train_Y)
        result = np.matmul(train_X.T, result)
        result = np.multiply(2, result)
        partial_der = result
        return partial_der

    # a function to calculate partial derivative of intercept
    def partial_derivative_of_intercept(self, weight, intercept, trainn_X, trainn_Y):

        xw = np.matmul(trainn_X, weight)
        xwi = xw + intercept
        result = np.subtract(xwi, trainn_Y)
        result = np.multiply(2, result)
        sum_of_elements = np.sum(result)
        partial_der_intercept = sum_of_elements
        return partial_der_intercept

    # a function to calculate squared loss for a given weight vector and intercept
    def calculate_loss(self, weight, trainnn_X, trainnn_Y, intercept):
        los = np.matmul(trainnn_X, weight)
        los = los + intercept
        los = np.subtract(los, trainnn_Y)
        los = np.matmul(los.T, los)
        return los

    def fit(self, x, y):

        trainx = x
        trainy = y
        weight_of_vector = len(trainx.columns)

        X = np.asarray(trainx)
        Y = np.asarray([trainy,])
        Y = Y.T

        # creating a random weight vector
        weight = []
        for w in range(0, weight_of_vector):

```

```
rand = np.random.normal(0,0.0001)
rand = np.asarray([rand,])
weight.append(rand)

weight = np.asarray(weight)

r = 1e-5 # learning rate
b = np.random.normal(0,10)
count = 0 # count to calculate number of iterations

# calculating squared loss for the random weights and intercept values.
loss = self.calculate_loss(weight,X,Y,b)

while True :
    Y= pd.DataFrame(np.array([trainy,]).T,columns=[ 'PRICE' ])
    X= pd.DataFrame(np.array(trainx))
    dataset = pd.concat([X,Y],axis = 1) # concatenating X,Y to sample the dataset
    dataset = dataset.sample(150) # sampling the dataset
    Y= pd.DataFrame(np.array(dataset[ 'PRICE' ]),columns=[ 'PRICE' ])
    X = dataset.drop(labels=[ 'PRICE' ],axis =1)
    X = pd.DataFrame(np.array(X))
    X = np.asarray(X)
    Y = np.asarray(Y)

    # calculating new weights
    term = self.partial_derivative_of_weights(weight,b,X,Y)
    term = np.multiply(r,term)
    term = np.subtract(weight,term)
    new_weight = term

    # calculating new intercept
    term2 = self.partial_derivative_of_intercept(weight,b,X,Y)
    term2 = np.multiply(r,term2)
    term2 = np.subtract(b,term2)
    new_intercept = term2

    # calculating new squared loss for the updated weights and intercept values.
    new_loss = self.calculate_loss(new_weight,X,Y,new_intercept)

    # threshold to stop the iterations
    threshold = np.absolute(loss - new_loss)
    loss = new_loss

    count = count + 1 # to count number of iterations
```

```
    r = r
    b = new_intercept
    weight = new_weight

    if threshold < 1.01:
        r = r/1.4 # to take smaller steps

    if threshold < 0.1 :
        self.intercept_ = new_intercept
        self.coef_ = new_weight
        self.n_iter_ = count
        break

    def predict(self,testing_data):
        pred_y = np.matmul(testing_data,self.coef_)
        pred_y = pred_y + self.intercept_
        return pred_y

sgd = Stochastic_Gradient_Descent()
sgd.fit(X_train,Y_train)
pred_sgd = sgd.predict(X_test)
```

Gradient Descent Implementation

```

In [355]: import numpy as np

class Gradient_Descent:

    # a function to calculate partial derivative of weights
    def partial_derivative_of_weights(self, weight, intercept, train_X, train_Y):
        xw = np.matmul(train_X, weight)
        xwi = xw + intercept
        result = np.subtract(xwi, train_Y)
        result = np.matmul(train_X.T, result)
        result = np.multiply(2, result)
        partial_der = result
        return partial_der

    # a function to calculate partial derivative of intercept
    def partial_derivative_of_intercept(self, weight, intercept, trainn_X, trainn_Y):

        xw = np.matmul(trainn_X, weight)
        xwi = xw + intercept
        result = np.subtract(xwi, trainn_Y)
        result = np.multiply(2, result)
        sum_of_elements = np.sum(result)
        partial_der_intercept = sum_of_elements
        return partial_der_intercept

    # a function to calculate squared loss for a given weight vector and intercept
    def calculate_loss(self, weight, trainnn_X, trainnn_Y, intercept):
        los = np.matmul(trainnn_X, weight)
        los = los + intercept
        los = np.subtract(los, trainnn_Y)
        los = np.matmul(los.T, los)
        return los

    def fit(self, x, y):

        trainx = x
        trainy = y
        weight_of_vector = len(trainx.columns)

        X = np.asarray(trainx)
        Y = np.asarray([trainy,])
        Y = Y.T

        # creating a random weight vector
        weight = []
        for w in range(0, weight_of_vector):

```

```
rand = np.random.normal(0,0.0001)
rand = np.asarray([rand,])
weight.append(rand)

weight = np.asarray(weight)

r = 1e-5 # learning rate
b = np.random.normal(0,10)
count = 0 # count to calculate number of iterations

# calculating squared loss for the random weights and intercept values.
loss = self.calculate_loss(weight,X,Y,b)

for hj in range(0,10000) :

    # calculating new weights
    term = self.partial_derivative_of_weights(weight,b,X,Y)
    term = np.multiply(r,term)
    term = np.subtract(weight,term)
    new_weight = term

    # calculating new intercept
    term2 = self.partial_derivative_of_intercept(weight,b,X,Y)
    term2 = np.multiply(r,term2)
    term2 = np.subtract(b,term2)
    new_intercept = term2

    # calculating new squared loss for the updated weights and intercept values.
    new_loss = self.calculate_loss(new_weight,X,Y,new_intercept)

    # threshold to stop the iterations
    threshold = np.absolute(loss - new_loss)
    loss = new_loss

    count = count + 1 # to count number of iterations

    r = r
    b = new_intercept
    weight = new_weight

    if threshold < 0.001 :
        self.intercept_ = new_intercept
        self.coef_ = new_weight
        self.n_iter_ = count
        break
```

```
def predict(self,testing_data):  
    pred_y = np.matmul(testing_data,self.coef_)  
    pred_y = pred_y + self.intercept_  
    return pred_y  
  
gd = Gradient_Descent()  
gd.fit(X_train,Y_train)  
pred_gd = gd.predict(X_test)
```

Similarity of all the three implementations


```
In [356]: import matplotlib.pyplot as plt



plt.figure(1)

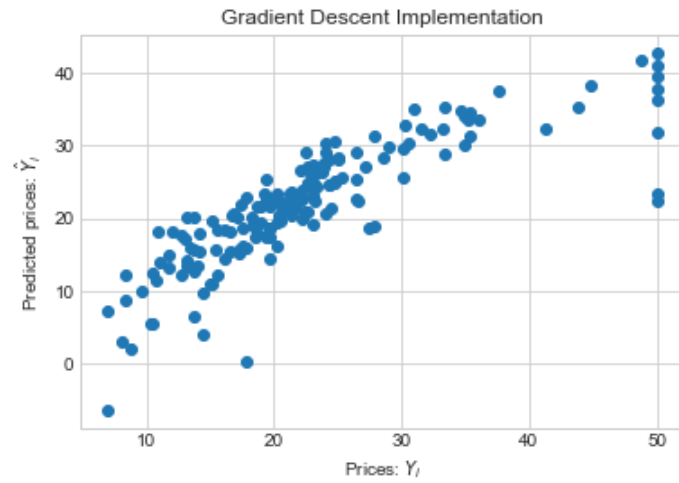
plt.scatter(Y_test, pred_sklern)
plt.xlabel("Prices: $Y_i$")
plt.ylabel("Predicted prices: $\hat{Y}_i$")
plt.title("Scikit-learn Implementation ")
plt.show()

plt.figure(2)
plt.scatter(Y_test, pred_sgd)
plt.xlabel("Prices: $Y_i$")
plt.ylabel("Predicted prices: $\hat{Y}_i$")
plt.title(" Stochastic Gradient Descent Implementation")
plt.show()

plt.figure(3)
plt.scatter(Y_test, pred_gd)
plt.xlabel("Prices: $Y_i$")
plt.ylabel("Predicted prices: $\hat{Y}_i$")
plt.title("Gradient Descent Implementation")

plt.show()
```





observation:

- * All the three scatter plots looks almost same.

```
In [357]: import seaborn as sns
import numpy as np

plt.figure(1,figsize=[15,10])


delta_y = Y_test - pred_sklern;
sns.set_style('whitegrid')
sns.kdeplot(np.array(delta_y), bw=0.5,label='Scikit-learn Implementation')

pred_sgd_ = pred_y.transpose()
pred_sgd_ = pred_sgd_[0]

delta_y = Y_test - pred_sgd_;
sns.set_style('whitegrid')
sns.kdeplot(np.array(delta_y), bw=0.5,label='Stochastic Gradient Descent Implementation')

pred_gd_ = pred_y.transpose()
pred_gd_ = pred_gd_[0]

delta_y = Y_test - pred_gd_;
sns.set_style('whitegrid')
sns.kdeplot(np.array(delta_y), bw=0.5,label='Gradient Descent Implementation')
plt.legend()
plt.show()
```



observation:

* Test data error distribution for all the three implementations are almost same.

```
In [358]: import seaborn as sns
import numpy as np

plt.figure(1,figsize=[15,10])

sns.set_style('whitegrid')


sns.kdeplot(np.array(pred_sklern), bw=0.5,label='Scikit-learn Implementation')

sns.kdeplot(np.array(pred_sgd_), bw=0.5,label='Stochastic Gradient Descent Implementation')

sns.kdeplot(np.array(pred_gd_), bw=0.5,label='Gradient Descent Implementation')

plt.legend()

plt.show()
```



observation:

* predicted price pdf's for all the three implementations are almost same.

```
In [361]: from prettytable import PrettyTable
from sklearn.metrics import mean_squared_error
from sklearn.metrics import r2_score

sklearn_mse = mean_squared_error(Y_test,pred_sklearn)
sgd_mse = mean_squared_error(Y_test,pred_sgd)
gd_mse = mean_squared_error(Y_test,pred_gd)

sklearn_r2 = r2_score(Y_test,pred_sklearn)
sgd_r2 = r2_score(Y_test,pred_sgd)
gd_r2 = r2_score(Y_test,pred_gd)

x = PrettyTable()

x.field_names = [ " ", "Scikit-learn implementation", " Gradient Descent", "Stochastic Gradient Descent", ]

x.add_row([ "Mean Squared Error\n", sklearn_mse, gd_mse, sgd_mse])
x.add_row([ "R2_Score\n", sklearn_r2, gd_r2, sgd_r2])

x.add_row([ "Intercept value\n", lm.intercept_, gd.intercept_, sgd.intercept_])

features = boston.feature_names
sklearn_weights = lm.coef_
sgd_weights = sgd.coef_
gd_weights = gd.coef_

for i in range(0,13):
    strr = str(features[i]) + " feature's weight\n"
    x.add_row([strr,np.round(sklearn_weights[i],decimals=3),np.round(gd_weights[i][0],decimals=3),np.round(sgd_weights[i][0],decimals=3)])

print(x)
print( " * Batch size of Stochastic Gradient Descent is 150")
```


	Scikit-learn implementation	Gradient Descent	Stochastic Gradient Descent
Mean Squared Error	28.541367275618263	28.508234880871253	29.11272209788439
R2_Score	0.6955388005506418	0.6958922359186075	0.6894439497742768
Intercept value	22.537168141592925	22.537168141592662	22.51753125430693
CRIM feature's weight	-1.312	-1.303	-1.107
ZN feature's weight	0.862	0.847	0.567
INDUS feature's weight	-0.167	-0.22	-0.54
CHAS feature's weight	0.19	0.197	0.241
NOX feature's weight	-1.487	-1.483	-0.918
RM feature's weight	2.791	2.801	3.01
AGE feature's weight	-0.327	-0.339	-0.411
DIS feature's weight	-2.772	-2.783	-2.313
RAD feature's weight	2.976	2.802	1.231
TAX feature's weight	-2.273	-2.08	-0.649
PTRATIO feature's weight	-2.134	-2.129	-2.031
B feature's weight	1.058	1.056	0.999
LSTAT feature's weight	-3.335	-3.328	-3.281
* Batch size of Stochastic Gradient Descent is 150			

-----THE END-----